NeuroLens

Watch a brain watch your video.

PREDICTED BRAIN RESPONSE TO VIDEO, SECOND BY SECOND

CS5224 CLOUD COMPUTING · FINAL PRESENTATION · DYLAN CHAN · XIN WEI CHONG · KELVIN LOW · JOSHUA LUM · BRANDON TAN

Video has never been easier to make.

So more of it is made, and all of it competes for the same attention.

Easier to make

New tools

Sora · Higgsfield · CapCut

More being made

A flood of video

AI dramas · ads · user-generated content

Harder to stand out

Marketing matters more

A market flooded with new ideas and products

You only find out after you post.

Until it’s live, there’s no way to know how viewers will react.



A FLOP MEANS THE MONEY AND EFFORT BEHIND IT ARE GONE

What if you could predict it?

NeuroLens simulates an fMRI scan of a brain watching your video.

PREDICTED RESPONSE (ILLUSTRATIVE)



WHERE

Which parts perform well

WHEN

Where viewers drop off

WHY

What the brain is doing

KNOW WHAT TO CHANGE BEFORE YOU POST · LESS EFFORT, MONEY AND TIME

What already exists: measured, filmed or guessed.

Physical fMRI study university and research labs	$10,000 – $15,000	2–4 weeks
Facial coding Realeyes, iMotions webcam studies	$850 – $4,000	3–5 days
AI attention maps Neurons, Dragonfly AI	$500 – $1,500 / mo	sub-minute
NeuroLens predicted brain response	$1.00	≈ 7 min

OUR DIFFERENCE

A brain response, not a gaze or a face

Per second, with a no-sound pass

Compare any two cuts directly

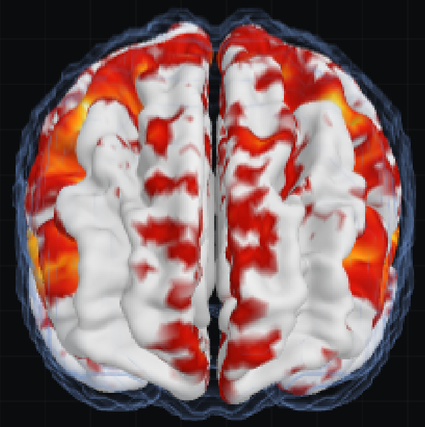
Under the hood: TRIBE v2.

Frames
V-JEPA 2 watches

Sound
Wav2Vec-BERT listens

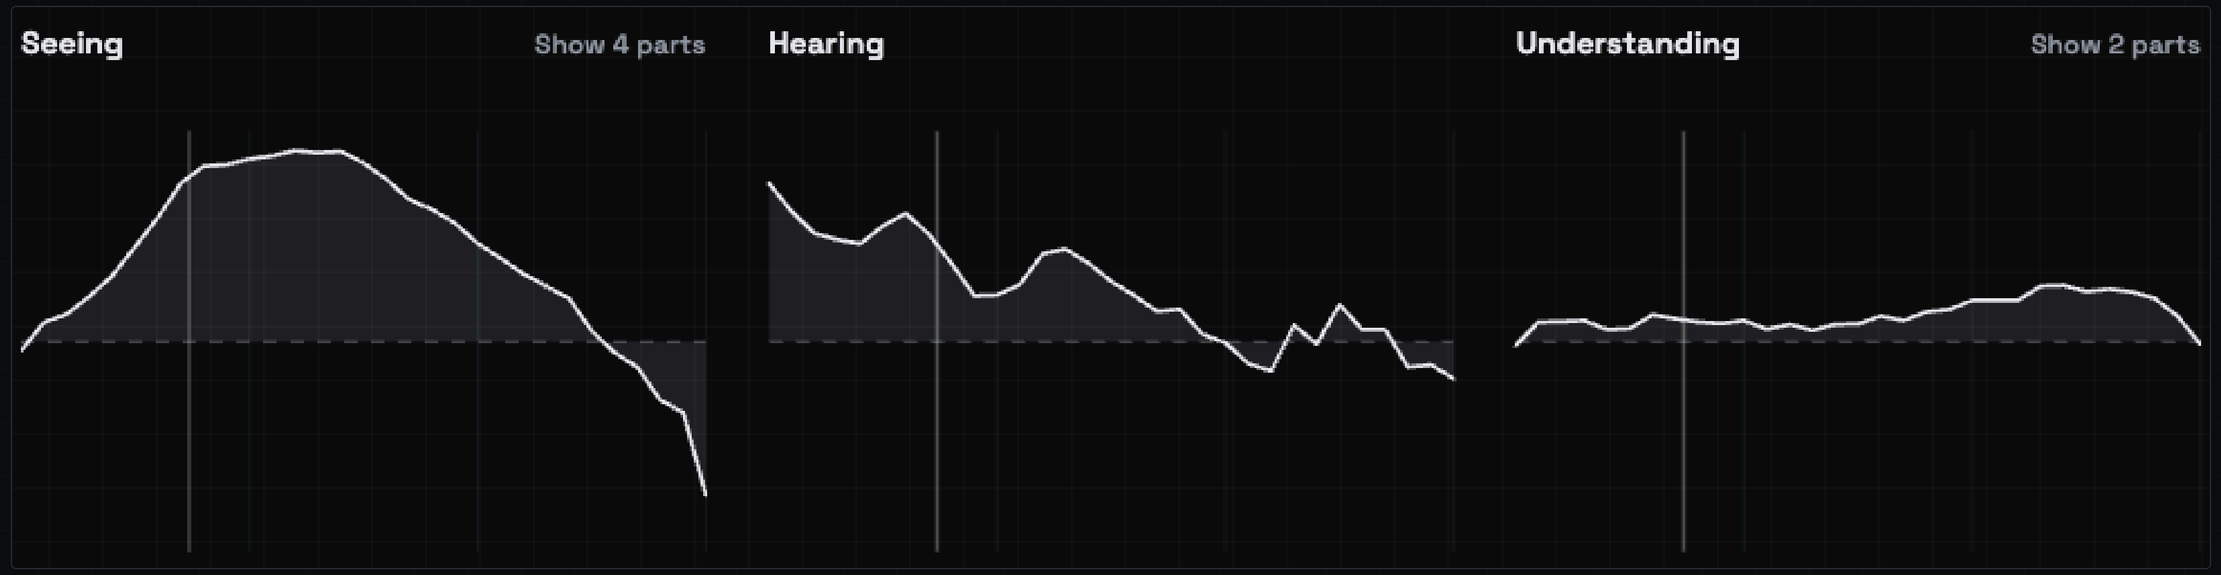
Words
WhisperX + LLaMA 3.2 read

TRIBE v2
1,000+ hours of fMRI
from 720 people



20,484 cortex points, every second

Seven signals, three questions.



SEEING

faces · bodies · scenes · visual

HEARING

sound

UNDERSTANDING

speech · social

BELOW TYPICAL

0 = TYPICAL FILM & TV VIEWING

ABOVE TYPICAL

Pay as you go. Credit, not cards.

$1.00

per started 30 s of video

NEW USERS: $10 CREDIT = TEN 30 S VIDEOS

WHY PREPAID CREDIT: CARD FEE AS A SHARE OF THE CHARGE

A $1.00 video, paid by card

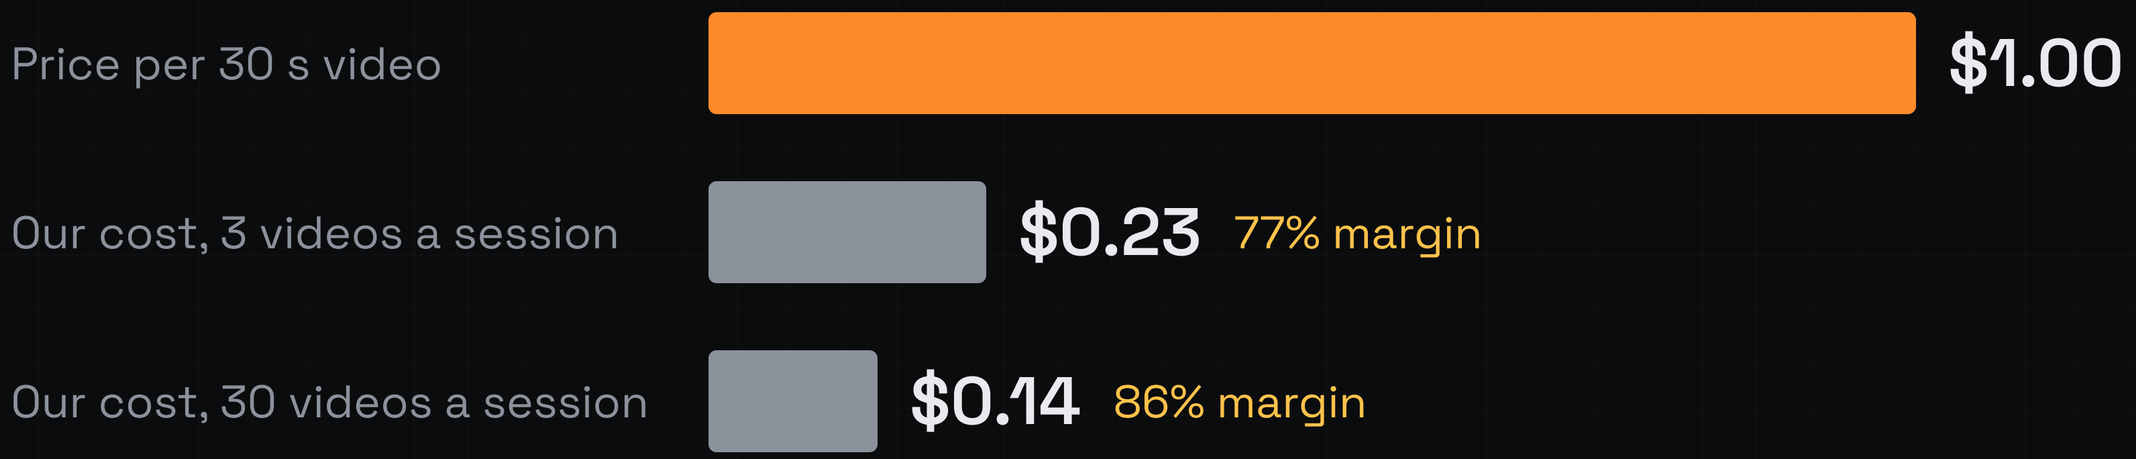
≈ 33%

A $20 credit pack

≈ 4%

COURSE VERSION: AN ADMINISTRATOR GRANTS CREDIT · STRIPE CHECKOUT BELONGS TO THE PRODUCTION DESIGN

What one analysis costs us.



GPU 393 S AT $1.22/H (L4) · START-UP AND IDLE SHARED ACROSS A SESSION · FIXED COST $5 A MONTH

Why cloud, not a workstation?

A GPU WORKSTATION

$159 / mo

about $5,000 up front over 3 years, plus idle power, before a single video is analysed

AD TESTING COMES IN BURSTS

80%+ idle

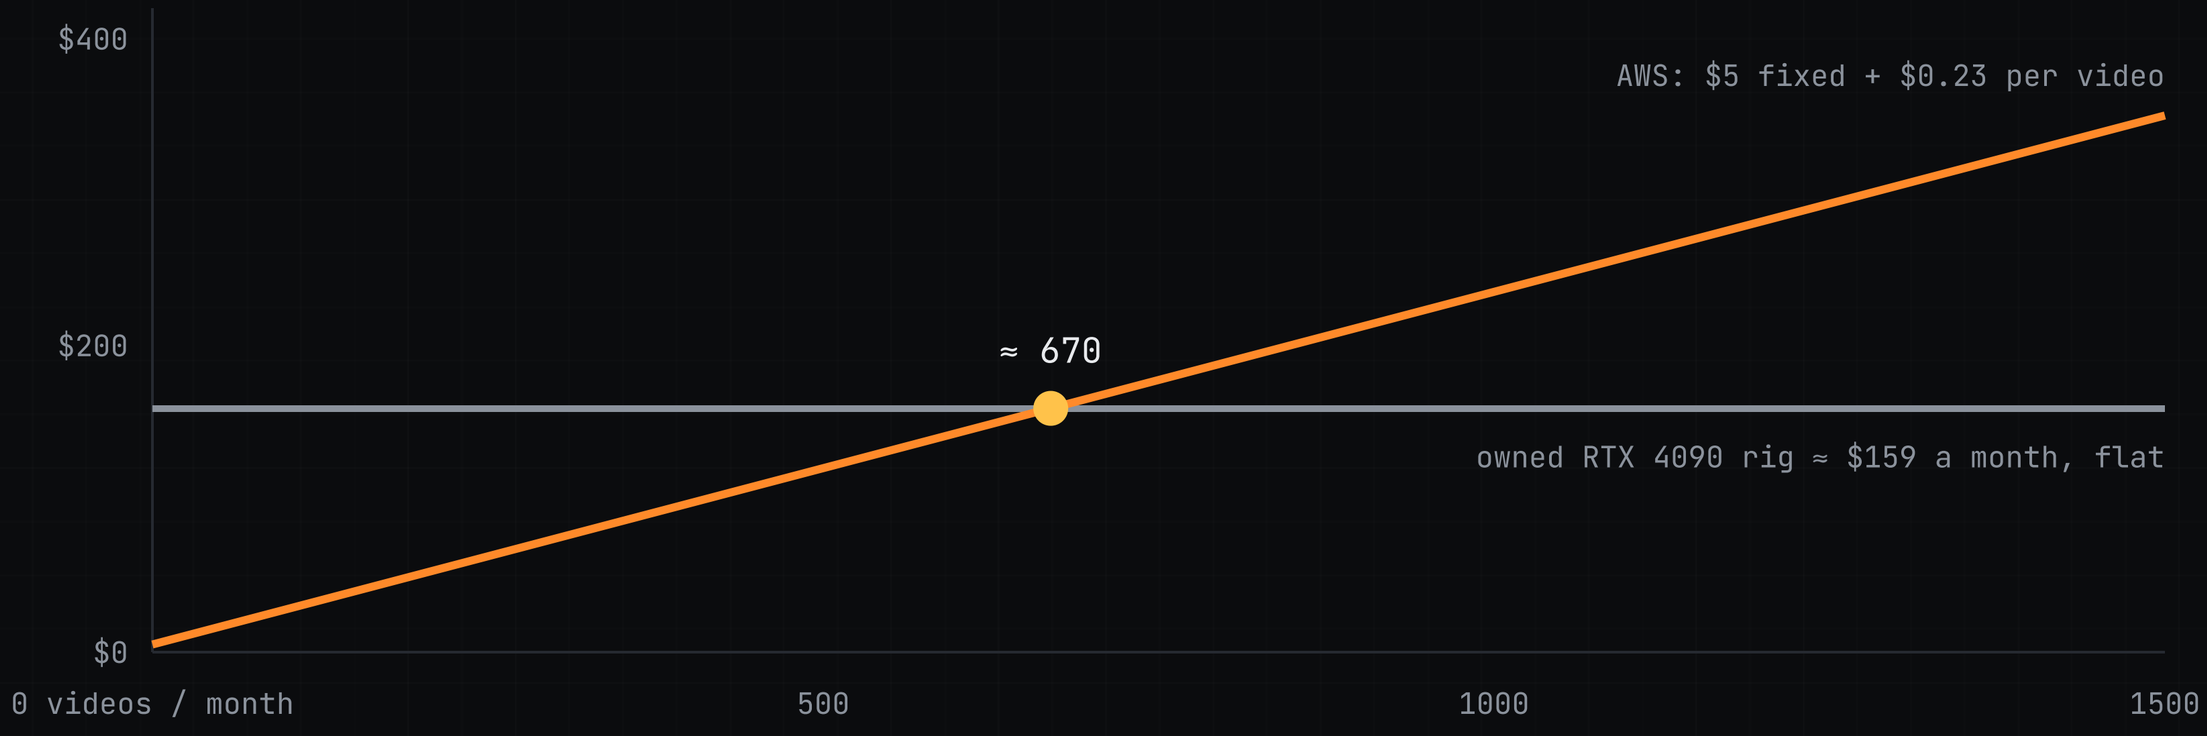
an agency may run 20 jobs in an editing sprint, then nothing for days

A CLOUD GPU, WHEN IDLE

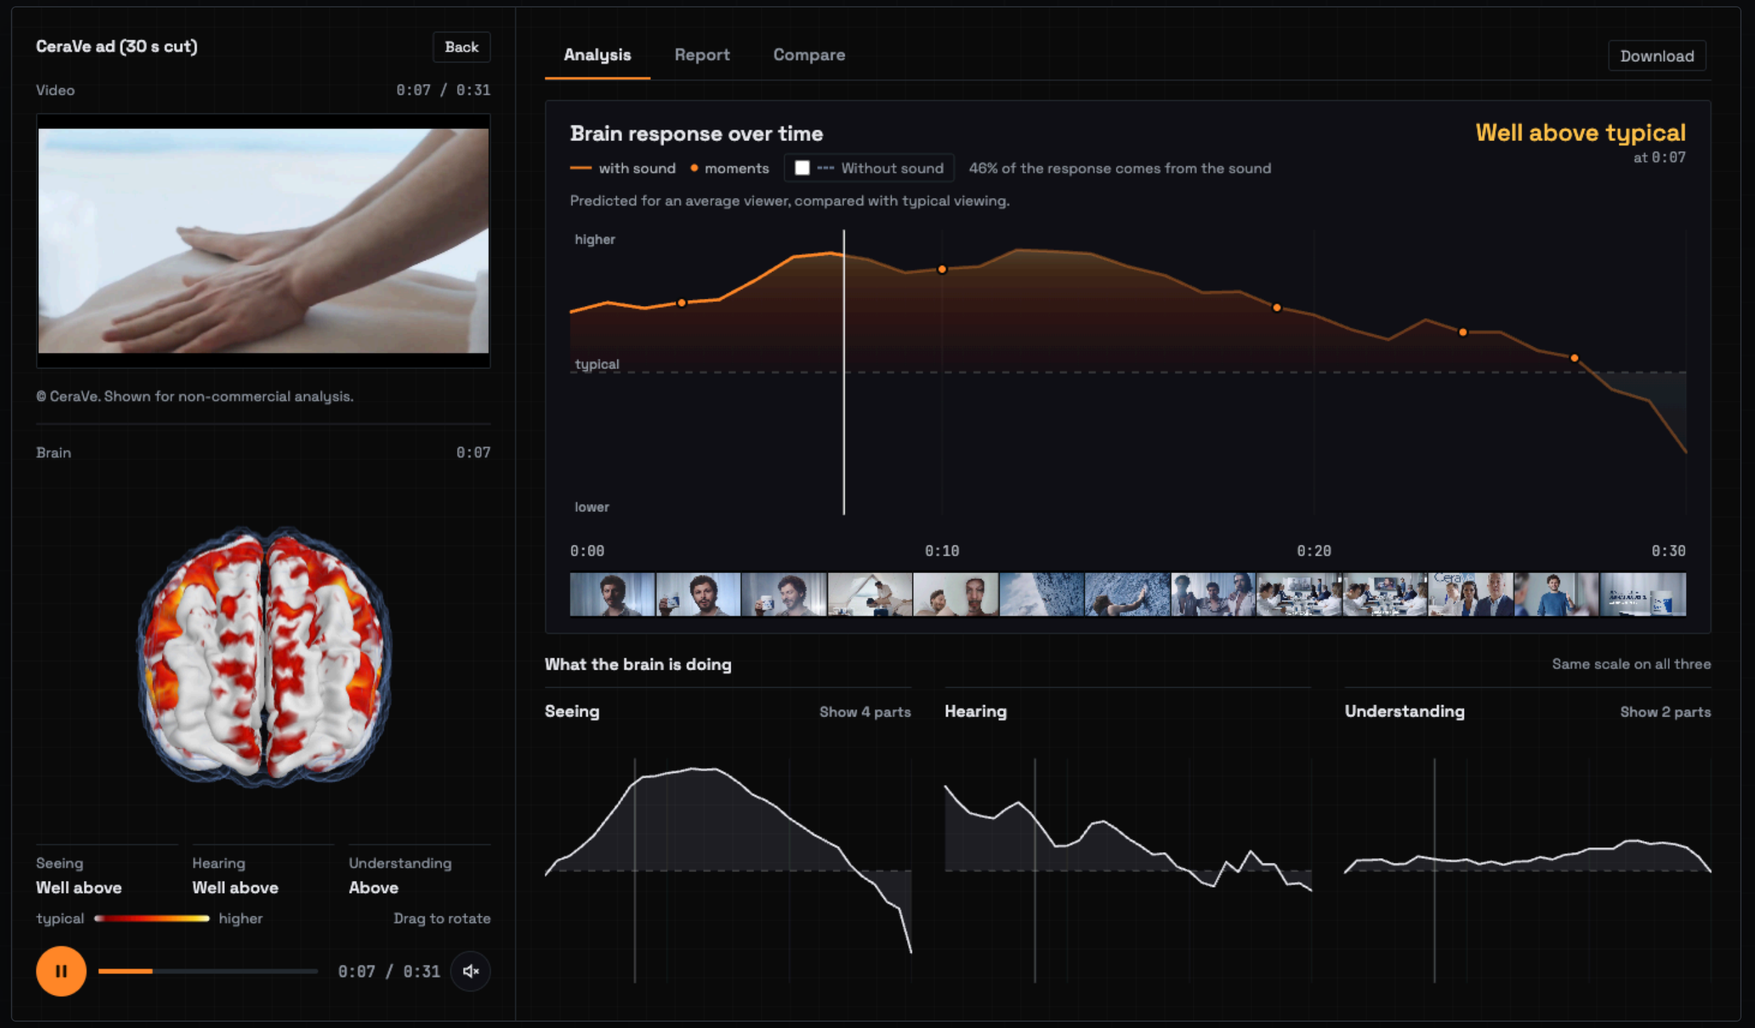
$5 / mo

it starts when a video is waiting and stops after 7 idle minutes. Always on, it would cost about $890 a month

Cloud is cheaper below about 670 videos a month.



LIVE Live demo: sign in, upload, explore



1. Upload: price first
2. Peak, dip, no-sound toggle
3. Seeing · hearing · understanding
4. AI report, compare, bad file

CLICK THE PICTURE TO OPEN THE LIVE SITE · DF6GIRQYYKYFA.CLOUDFRONT.NET

Can we trust it?

SAME ANSWER ON TWO MACHINES

0.005

Largest difference across 11 series, same clip on our RTX 4090 and on AWS. Every series correlates at 1.000.

FAILED JOBS ARE REFUNDED

77 / 80

Automated tests: a failed or killed job refunded exactly once. The other 3 stopped on a GPU-library error in setup.

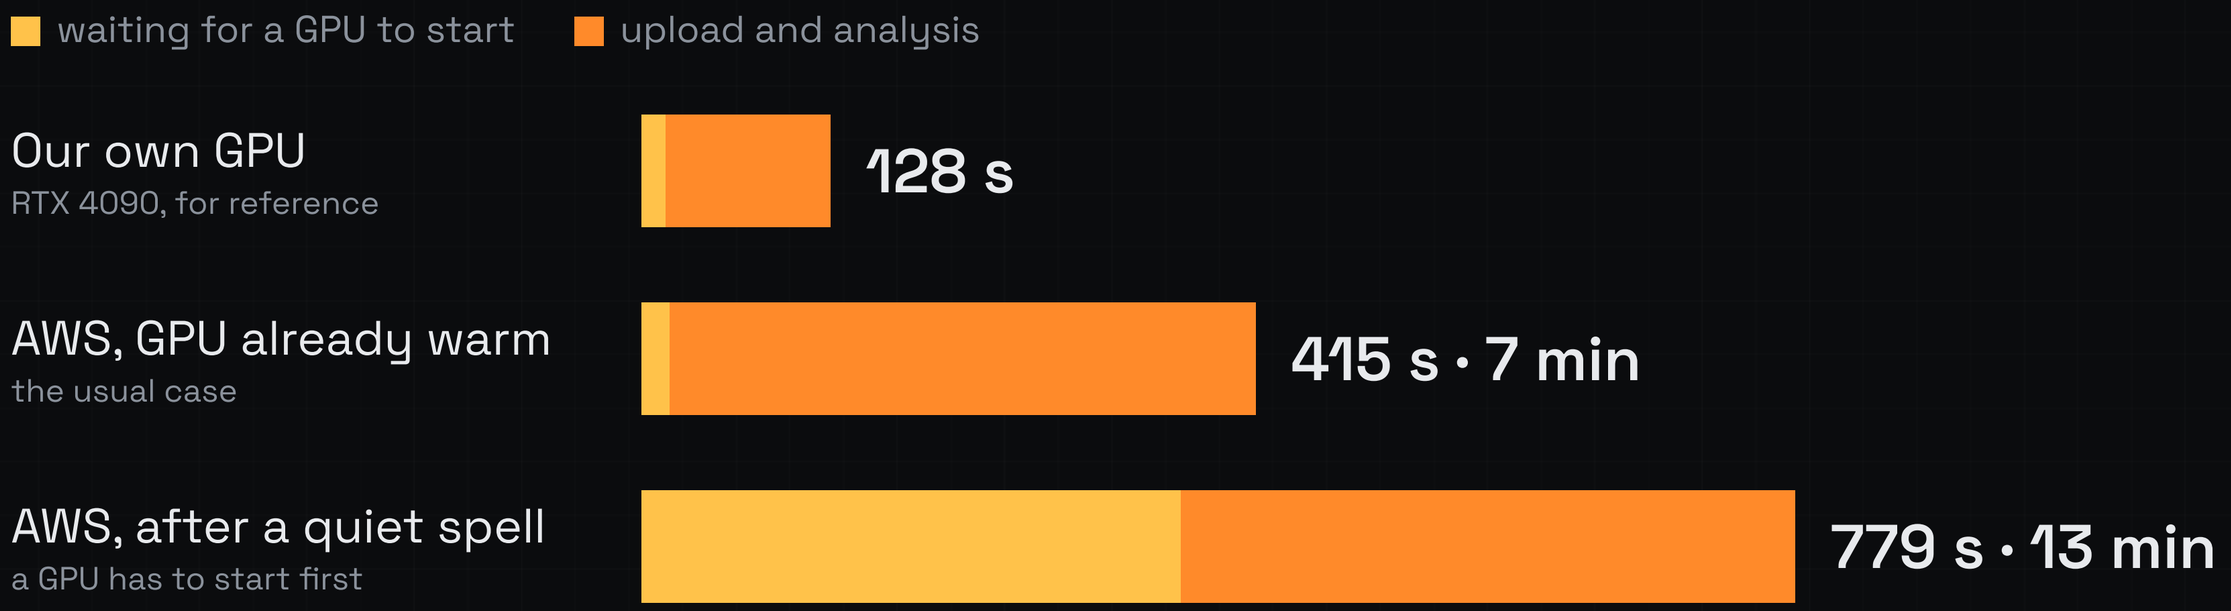
NOT TESTED

Views, sales

Whether a higher response means more views or sales. We make no such claim.

TRIBE V2'S ACCURACY IS META'S RESULT (NEW VIDEOS, TASKS AND PEOPLE); WE DID NOT RE-TEST IT AGAINST REAL SCANS

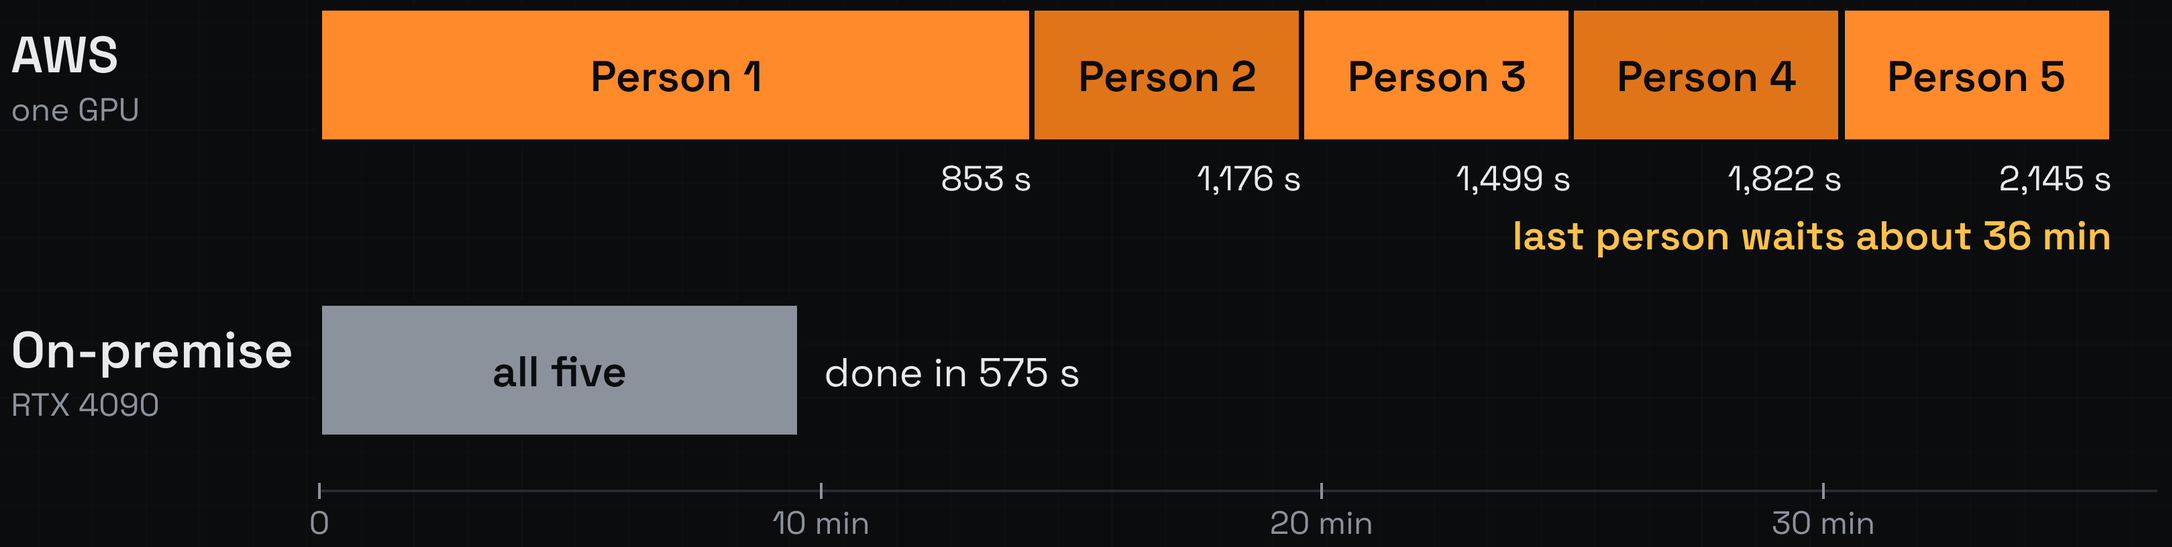
How long will a user wait?



SAME 30 S CLIP · GPU WORK PER VIDEO-SECOND: 3.7 S ON THE RTX 4090, 10.7 S ON AN A10G · COLD-START WAITS: 452 S AND 277 S

Five people at once: they wait in line.

Five uploads at the same moment. One GPU runs one job at a time.



NOTHING WAS DROPPED. THE AWS RUN GOT AN A10G, AND ITS FIRST JOB APPEARS TO HAVE RESTARTED ONCE.

Can people use it?

71.6 / 100

SUS SCORE · 11 PEOPLE · 68 IS AVERAGE · 80 IS VERY GOOD



6 scored 77.5 or more · 5 scored 47.5 to 55 · 5 of 11 reached 80

WHAT CONFUSED PEOPLE (2 OF 11)

A muted video looked like a phone fault

Fix: a sound control on each player

What do the brain regions mean?

Fix: a short note on each signal

8 OF 11 TESTERS WERE STUDENTS, NOT OUR TARGET CREATORS AND AGENCIES

Is it worth using?

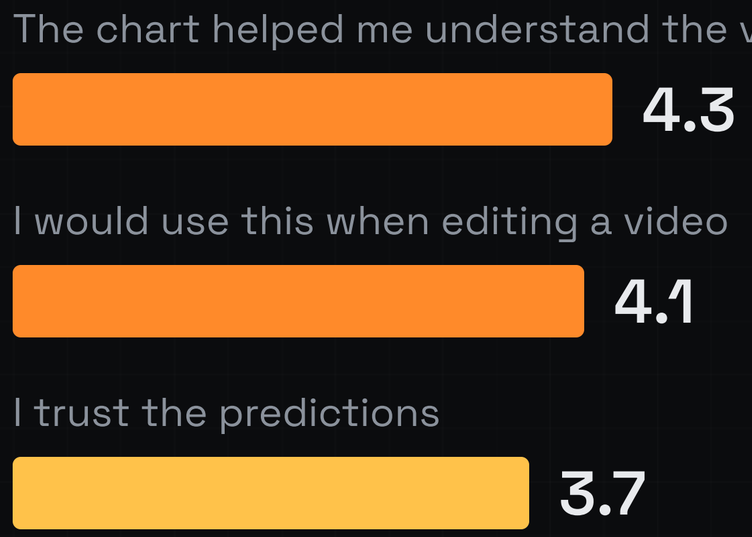
82%

said it showed them something they had not noticed (9 of 11)

64%

would pay $1.00 for a 30 s video (7 of 11); 2 more might

AVERAGE AGREEMENT, 1 TO 5



“It helps me improve my video easily and update the scenes to make it more engaging.”

MAIN DOUBT: ACCURACY. 3 OF THE 4 WHO WOULD NOT PAY, OR WERE UNSURE, SAID SO

What broke, and what we did.

FIXED

Words repeated in every chunk

Before: all words in every chunk



After: each chunk, its own words

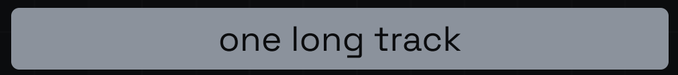


We replaced TRIBE v2's speech step.

FIXED

Long audio ran out of GPU memory

Example: 150 s of audio



Split into equal pieces ≤ 60 s

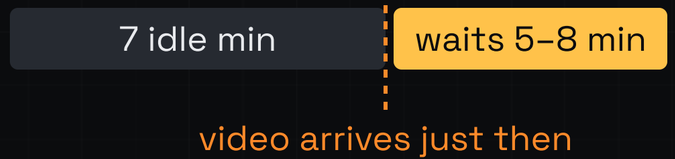


We split audio over 60 s into equal pieces.

KNOWN LIMIT

A video arrives as the GPU stops

The GPU stops after 7 idle min



The request restarts the GPU, so it waits a further 5–8 min.

What it claims, and what it doesn’t.

WHAT IT IS

- An average viewer’s brain response
- One scale for every video

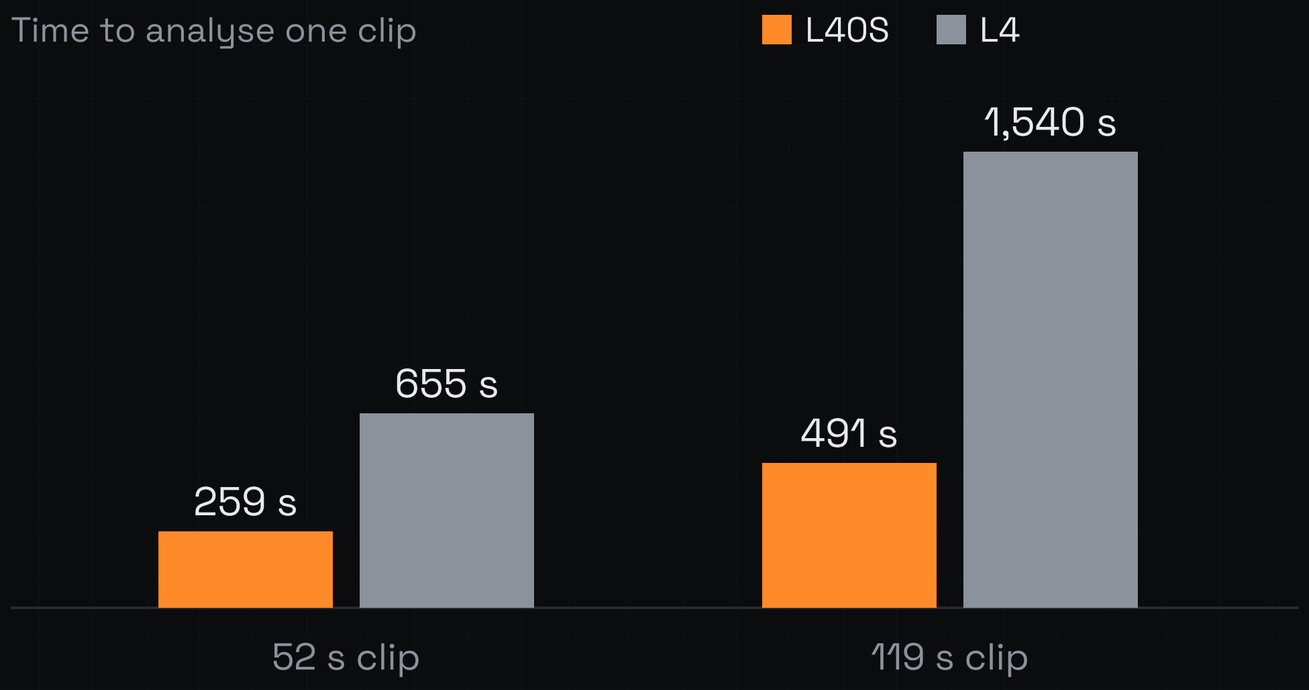
WHAT IT ISN'T

- A prediction of views or sales
- Instant: first video ≈ 6 min
- Tested on creators (8 of 11 were students)

ETHICS

- Non-commercial licence (CC BY-NC)
- No one scanned or profiled
- Anonymous, consented survey

Where next: scaling up.



26 / 26

Spot launches failed for lack of capacity

19.9 GB

peak GPU memory: a 24 GB card fits

An EC2 GPU pool: $0.10 – $0.20 a video, about $70 more a month. It pays off above about 1,000 videos a month. Estimates, not built.

A $10,000 brain study, in 7 minutes for $1.

$1.00
per 30 s video

≈ 7 min
to a result, GPU warm

670
videos a month: break-even vs owning

≈ 6 min
first video after a quiet spell

Next

Licence

CC BY-NC: a paid launch needs one.

Next

Scale

A GPU pool that grows with the queue.

Next

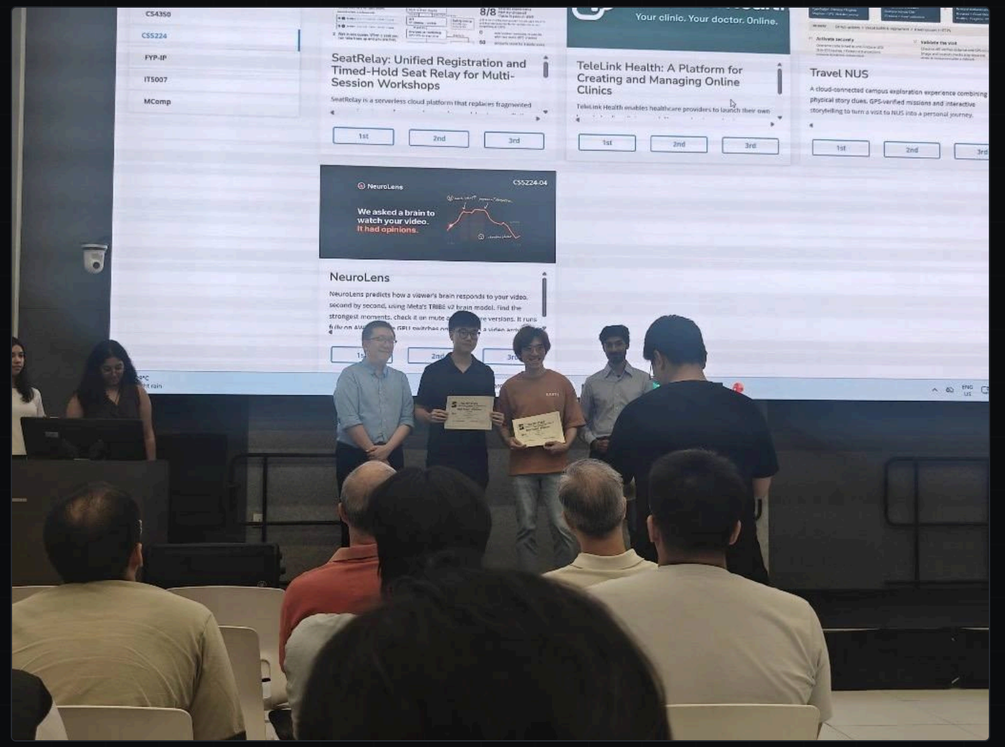
Polish

A sound control on each player, a note on each signal.

Thank you. Questions?

df6girqyykyfa.cloudfront.net

CODE · GITHUB.COM/KELVINLOWHM86/NEUROLENS



Thanks for supporting us at STePS!